

Xiang LIU

Homepage: <https://xiangli-ml.github.io/>

Google Scholar: [VtK5lwUAAAAJ](#) (1106 citations)

Email: xliu886@connect.hkust-gz.edu.cn

Mobile: +86-15617865777

EDUCATION

-
- **The Hong Kong University of Science and Technology (Guangzhou)** **Guangzhou, CHINA**
*PhD in DSA Thrust, supervised by **Professor Xiaowen Chu*** *Sep 2023 - Aug 2027*
 - **New York University** **NY, USA**
*Visiting PhD in CDS, supervised by **Professor Eunsol Choi*** *July 2025 - Jan 2026*
 - **The University of Hong Kong** **Hong Kong SAR, CHINA**
Master of Science(MSc) in Computer Science *Sep 2022 - Aug 2023*
 - **George Mason University** **VA, USA**
Bachelor of Science in Computer Science; GPA: 3.71/4.0 *Aug 2018 - Aug 2022*
Honors/Awards: Dean's List (2018-2020)

SELECTED RESEARCH

First / co-first author: 10 papers (ICML $\times 2$ · ICLR $\times 2$ · NeurIPS $\times 3$ · EMNLP Findings $\times 2$ · Workshop $\times 1$)

- **SIM: Benchmarking and Preserving High-Density Reasoning in KV Cache Compression** **ICML 2026**
***First Author** · KV Cache, Reasoning, Compression, Benchmark*
- **MDN: Parallelizing Stepwise Momentum for Delta Linear Attention** **ICML 2026**
***Co-first Author** · Linear Attention, Delta Rule, Parallelization, Momentum*
- **DiffAdapt: Difficulty-Adaptive Reasoning for Token-Efficient LLM Inference** **ICLR 2026**
***First Author** · Adaptive Inference, Reasoning LLM, Token Efficiency, Entropy Analysis*
- **Reasoning Language Model Inference Serving Unveiled: An Empirical Study** **ICLR 2026**
***Co-first Author** · LLM Serving, Reasoning System, Empirical Study, Memory Optimization*
- **ChunkKV: Semantic-preserving KV cache compression for efficient long-context LLM** **NeurIPS 2025**
***First Author** · KV Cache, Inference, GPU Memory, Long-Context*
- **Perovskite-LLM: Knowledge-Enhanced LLMs for Solar Cell Research** **EMNLP Findings 2025**
***Co-first Author** · Domain-Specific LLM, Knowledge Augmentation, Materials Science*
- **LISA: Layerwise Importance Sampling for Memory-Efficient LLM Fine-Tuning** **NeurIPS 2024**
***Co-first Author** · Parameter Efficient Fine-tuning, Layer-wise Optimization*
- **Should We Really Edit Language Models? On the Evaluation of Edited LLMs** **NeurIPS 2024**
***Co-first Author** · Model Editing, Benchmark, Model Robustness*
- **LongGenBench: Long-context Generation Benchmark** **EMNLP Findings 2024**
***First Author** · Long-context Generation, Long-context LLMs, Logical Coherence*

EXPERIENCE

-
- **Mind Lab** **Remote**
Research Intern *April 2026 - Present*
 - Post-trained **GLM-5.1** and **Qwen3.6 35B** for tool use and code generation; improved agent capability on *ViTaBench, PinchBench, Tau3Bench, SWE-Bench*, and *Terminal2Bench*.
 - Strengthened model harness capability for long-horizon tool calling and code execution.
 - Co-authored **δ -mem** and **MinT** preprints.
 - **NYU Center for Data Science** **NY, USA**
Visiting Student Supervisor: Prof. Eunsol Choi *July 2025 - Jan 2026*
 - Proposed **DiffAdapt**, a novel algorithm for efficient adaptive reasoning for LLMs (Accepted to ICLR 2026).
 - Conducted comprehensive entropy analysis on reasoning traces to identify "overthinking" phenomena.
 - **HKUST Statistics and Machine Learning Research Group** **HK, CHINA**
Research Intern Supervisor: Prof. Tong Zhang *Dec 2022 - Aug 2023*
 - Proposed **LISA**, a novel algorithm for efficient fine-tuning of LLMs, accepted at *NeurIPS 2024*.
 - Contributed to the LMFlow framework, enabling personalized LLM fine-tuning and deployment (LLaMA, Bloom, Vicuna).
 - Conducted research on Chain-of-Thought (CoT) methods to enhance logical reasoning in LLMs.
 - **Baidu Research Cognitive Computing Lab** **Beijing, CHINA**
Research Intern *Dec 2021 - June 2022*
 - Worked on dependency parsing using Open Information Annotation (OIA) to convert sentences into DAGs.

COMPETITIONS

- **Kaggle — Feedback Prize - Predicting Effective Arguments Competition** **Jun 2022 - Aug 2022**
 - *Team Leader — Silver Medal (Top 2%)*
 - Developed baseline code, designed data preprocessing strategies, and optimized model structures.
 - Implemented token classification replacing sequence classification, boosting rank and efficiency.

PROFESSIONAL SKILLS

- **Languages:** English, Mandarin
- **Computing Skills:** PyTorch, Git, Linux